

Azure Stream Analytics

Manjunath Suryanarayana



Unlocking Real-time Insights

Time to Insight is Critical


Reducing decision latency can unlock business value

Insights are Perishable

Window of opportunity for insights to be actionable

Ask Questions to Data in Motion

Can't wait for data to get to rest before running computation



Real-time Stream Processing

Simple Event Processing

- Filter
- Transform
- Enrich
- Split
- Route

Event Stream Processing

- [Simple event processing] +
- Aggregate
- Rules

Complex Event Processing

- [Event Stream Processing] +
- Pattern detection
- Time windows
- Joins & correlations



Scenarios



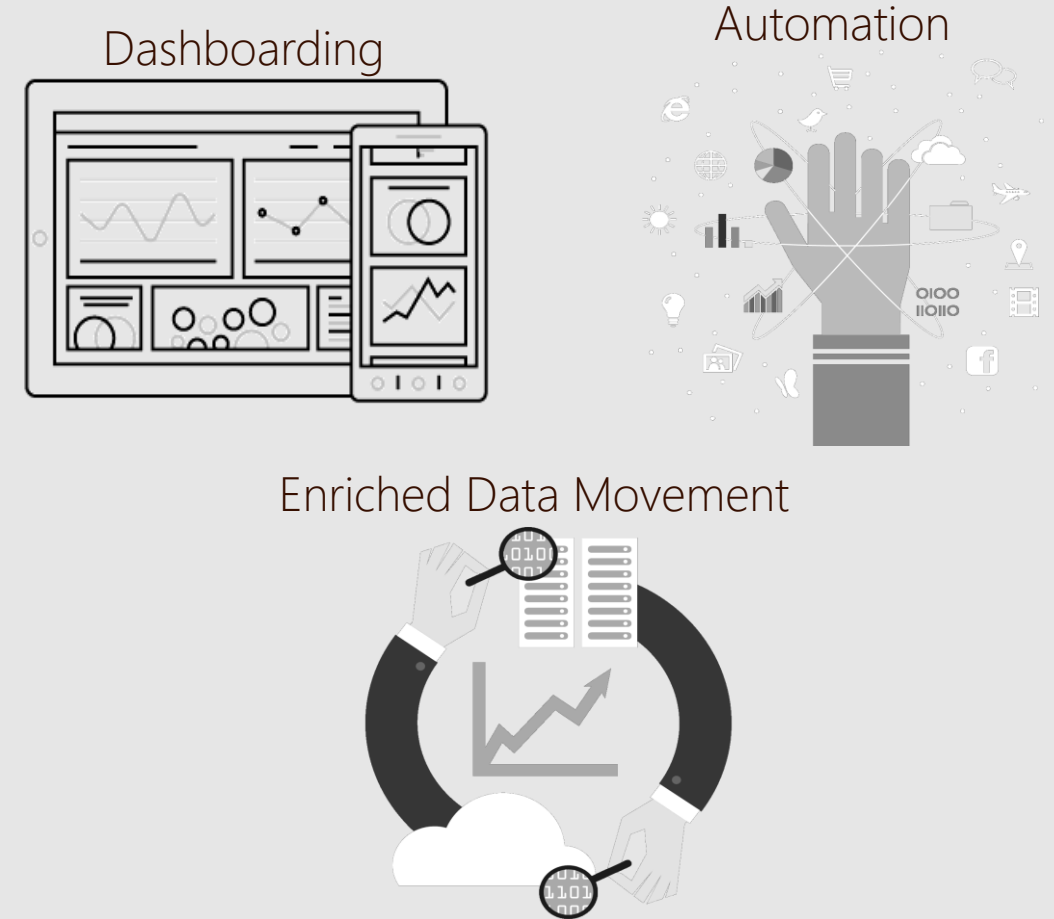
Scenario Types

Actions by Human Actors

- "See and seize" insights
- Live visualization
- Alerts and alarms
- Dynamic aggregation

Machine to Machine Interactions

- Data movement with enrichment
- Kick-off workflows for automation



Scenario Examples

Real-time Fraud Detection




Streaming ETL




Predictive Maintenance



Call Center Analytics



IT Infrastructure and Network Monitoring



Customer Behavior Prediction



Log Analytics



Real-time Cross Sell Offers



Fleet monitoring and Connected Cars




Real-time Patient Monitoring



Smart Grid



Real-time Marketing

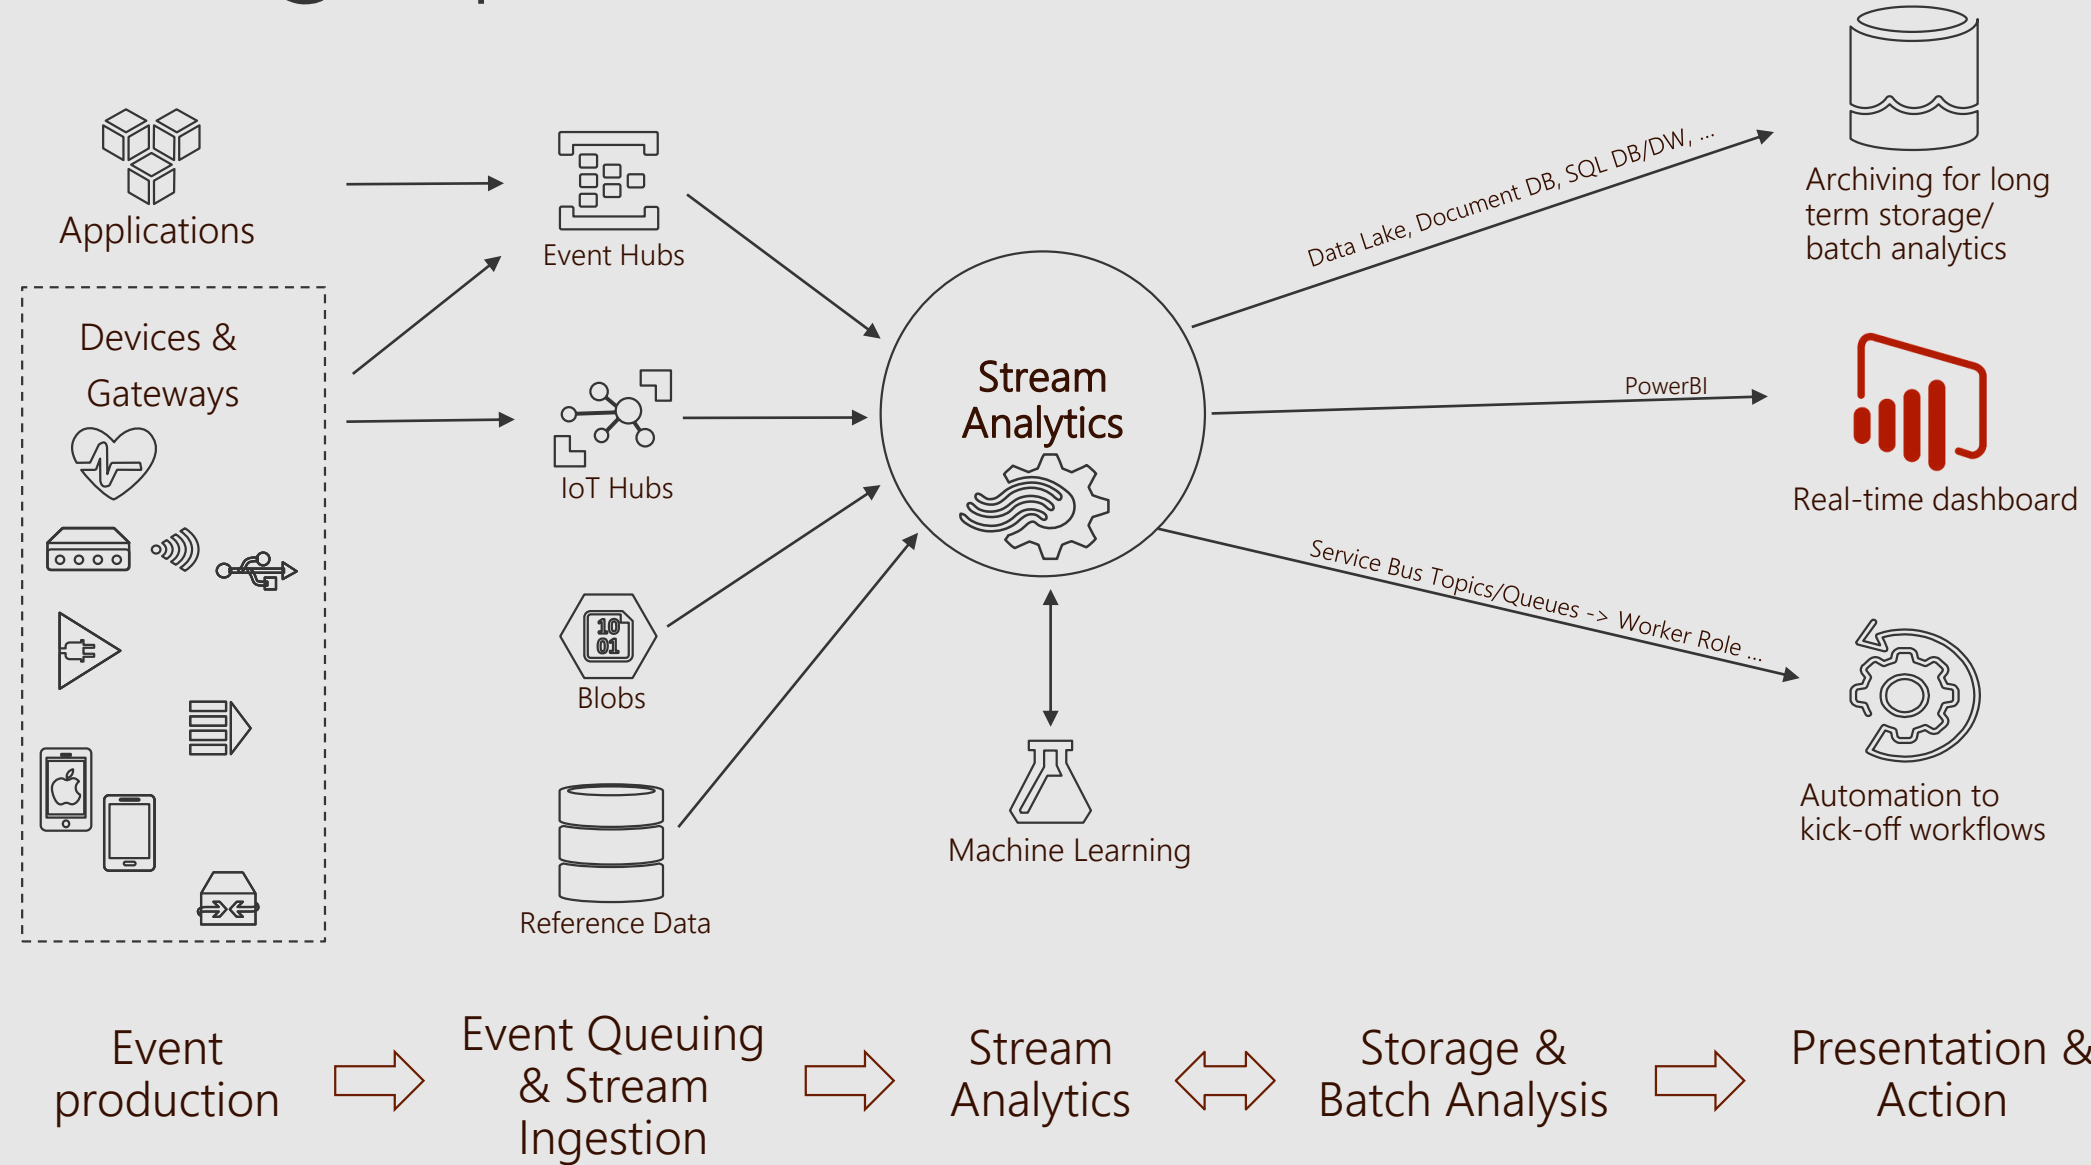


and many more...

Streaming Analytics Pipeline



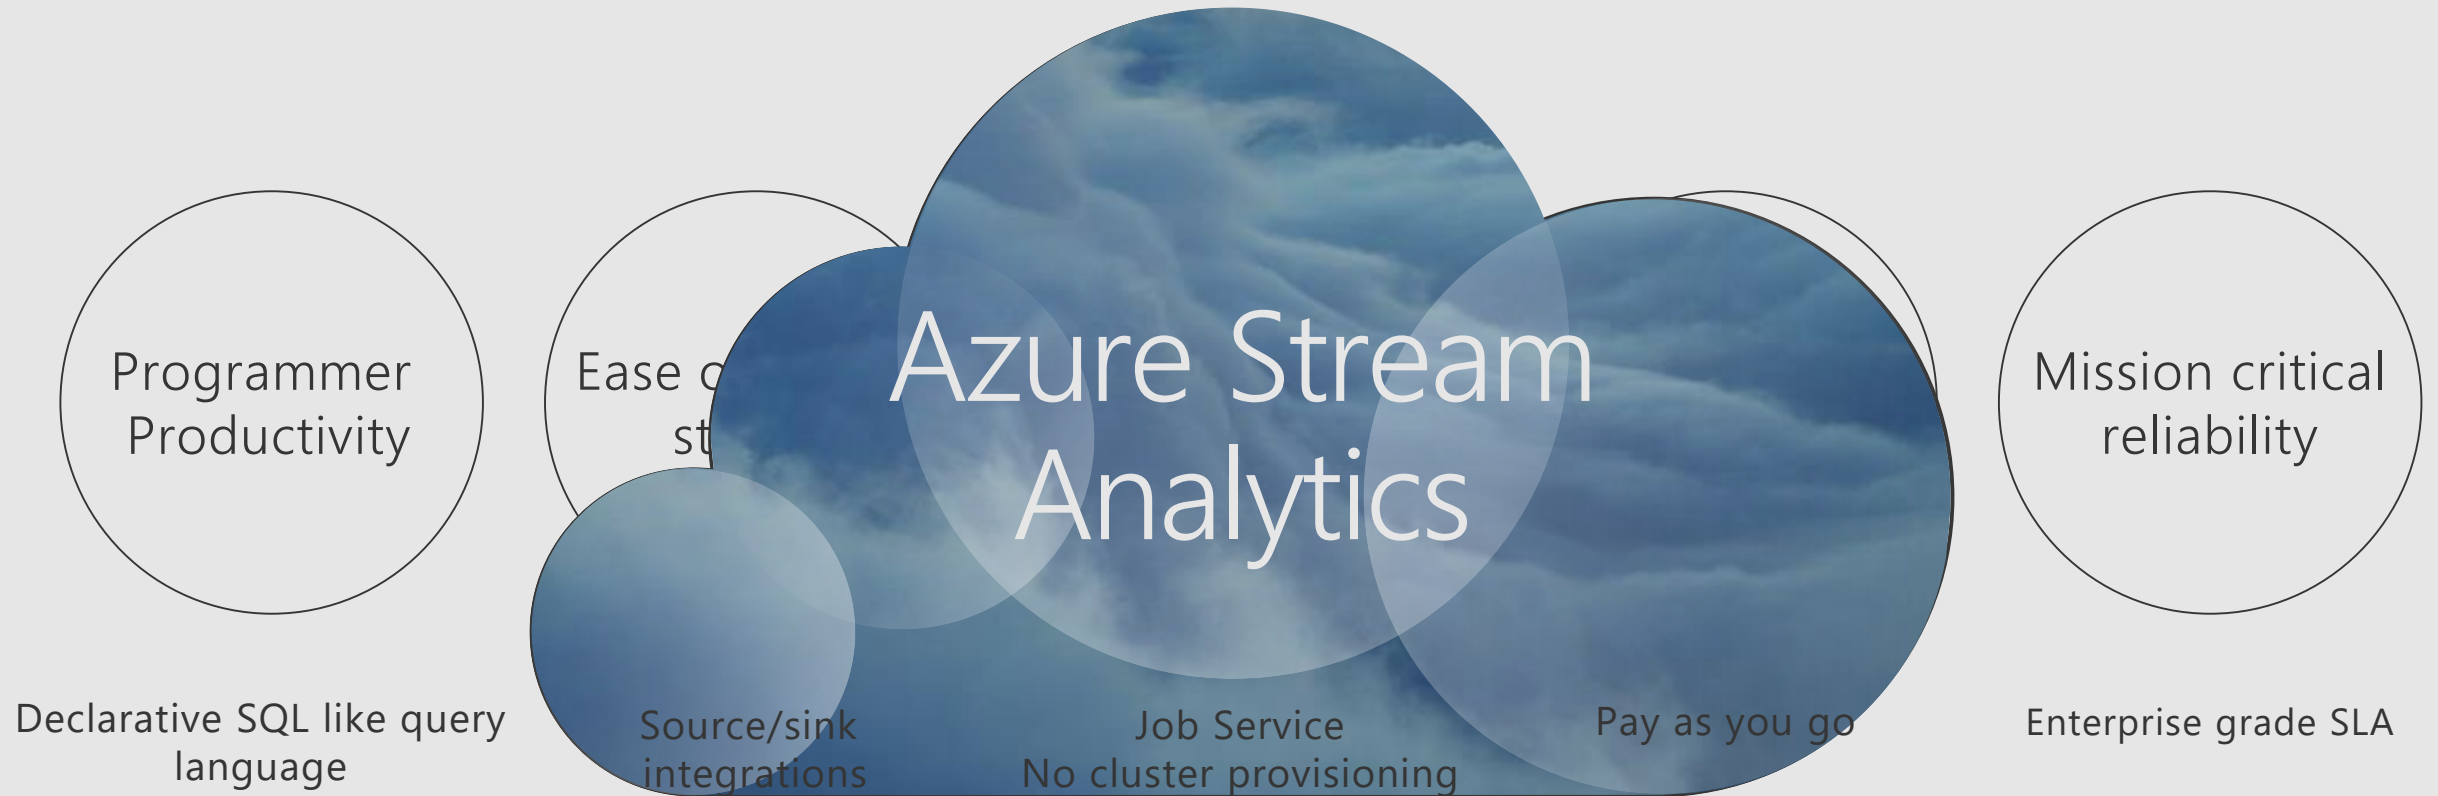
Streaming Pipeline



Differentiated Value Proposition



Azure Stream Analytics



Differentiators

Programmer Productivity

- Declarative SQL like language
- Built-in temporal semantics

Ease of Getting Started

- Integrations with sources, sinks, & ML
- Build real-time dashboards in minutes

Lowest Total Cost of Ownership(TCO)

- Fully managed service
- No cluster topology management required
- Seamless scalability
- Usage based pricing

1,915 lines of code with Apache Storm

```
@ApplicationAnnotation(name="WordCountDemo")
public class Application implements StreamingApplication
{
    protected String fileName =
        "com/datatorrent/demos/wordcount/samplefile.txt";
    private Locality locality = null;

    @Override public void populateDAG(DAG dag, Configuration
    conf)
    {
        locality = Locality.CONTAINER_LOCAL;
        WordCountInputOperator input =
            dag.addOperator("wordinput", new
            WordCountInputOperator());
        input.setFileName(fileName);
        UniqueCounter<String> wordCount =
            dag.addOperator("count", new
            UniqueCounter<String>());
        dag.addStream("wordinput-count", input.outputPort,
            wordCount.data).setLocality(locality);
        ConsoleOutputOperator consoleOperator =
            dag.addOperator("console", new
            ConsoleOutputOperator());
        dag.addStream("count-console", wordCount.count,
            consoleOperator.input);
    }
}
```

3 lines of SQL in Azure Stream Analytics

```
SELECT Avg(Purchase), ScoreTollId, Count(*)
FROM GameDataStream
GROUP BY TumblingWindows(5, Minute), Score
```

Stream Analytics Query Language (SAQL)

Declarative SQL like language to describe transformations

- Filters ("Where")
- Projections ("Select")
- Time-window and property-based aggregates ("Group By")
- Time-shifted joins (specifying time bounds within which the joining events must occur)
- and all combinations thereof

Data Manipulation

SELECT
FROM
WHERE
HAVING
GROUP BY
CASE WHEN THEN
ELSE
INNER/LEFT OUTER
JOIN
UNION
CROSS/OUTER APPLY
CAST INTO
ORDER BY ASC, DSC

Aggregation

SUM
COUNT
AVG
MIN
MAX
STDEV
STDEVP
VAR
VARP
TopOne

Date and Time

DateName
DatePart Day, Month, Year
DateDiff
DateTimeFromParts
DateAdd

Temporal

Lag
IsFirst
Last
CollectTop

Windowing Extensions

TumblingWindow
HoppingWindow
SlidingWindow

Scaling Extensions

WITH
PARTITION BY
OVER

String

Len
Concat
CharIndex
Substring
Lower, Upper
PatIndex

Mathematical

ABS
CEILING
EXP
FLOOR
POWER
SIGN
SQUARE
SQRT

Geospatial (preview)

CreatePoint
CreatePolygon
CreateLineString
ST_DISTANCE
ST_WITHIN
ST_OVERLAPS
ST_INTERSECTS

Mission Critical Reliability

Enterprise Grade SLA

At least three 9s of availability

Business Continuity During Failures


Automatic checkpoint-recovery

Fast restarts

Guaranteed Event Delivery

At-least-once event delivery semantics

No data loss



Global Availability Footprint

Currently available in 19 Azure regions including China and Germany

Current list includes:

Central US, East US, East US2, North Central US, South, Central US, West US, North Europe, West Europe, East Asia, Southeast Asia, Japan West, Japan East, Brazil South, Australia East, Central India.

In China: Stream Analytics is made available through a unique partnership between Microsoft and 21Vianet.

In Germany: Stream Analytics is available via a new data trustee model whereby customer data remains in Germany under control of T-Systems, a Deutsche Telekom company, acting as the German data trustee.



Satisfies Major Global Compliance Requirements

Current list includes:

ISO 27001

ISO 27018

SOC 1 Type 2

SOC 2 Type 2

SOC 3 Type 2

HIIPAA/HITECH

PCI DSS Level 1

European Union Model Clauses

China GB 18030



Capabilities



Stream Analytics Job

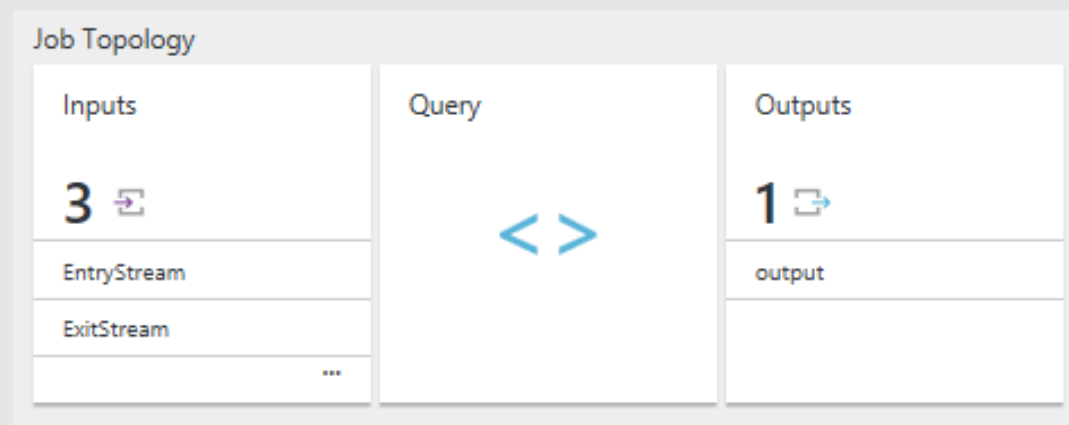
Users construct and deploy jobs to ASA

Job definition includes inputs, a query, and output

Inputs are from where the job reads the data stream

Query runs for perpetuity unless explicitly stopped and transforms the input stream

Output is where the job sends the job results to



Scenario – Tolling Station

Tolling stations have multiple booths (identified by TollId)

Each booth has two sensors: Entry and Exit that send out EntryStream and ExitStream respectively

EntryStream – Data stream from the Entry sensor data on vehicles entering toll booths

| TollId | EntryTime | License Plate | State | Make | Model | Type | Weight |
|--------|------------------------------|---------------|-------|--------|---------|------|--------|
| 1 | 2014-10-25T19:33:30.0000000Z | JNB 7001 | NY | Honda | CRV | 1 | 3010 |
| 1 | 2014-10-25T19:33:31.0000000Z | YXZ 1001 | NY | Toyota | Camry | 2 | 3020 |
| 3 | 2014-10-25T19:33:32.0000000Z | ABC 1004 | CT | Ford | Taurus | 2 | 3800 |
| 2 | 2014-10-25T19:33:33.0000000Z | XYZ 1003 | CT | Toyota | Corolla | 2 | 2900 |
| ... | ... | ... | ... | ... | ... | ... | ... |

ExitStream - Data stream from the Exit sensor on vehicles exiting toll booths

| TollId | ExitTime | LicensePlate |
|--------|------------------------------|--------------|
| 1 | 2014-10-25T19:33:40.0000000Z | JNB 7001 |
| 1 | 2014-10-25T19:33:41.0000000Z | YXZ 1001 |
| 3 | 2014-10-25T19:33:42.0000000Z | ABC 1004 |
| 2 | 2014-10-25T19:33:43.0000000Z | XYZ 1003 |
| ... | ... | ... |

ReferenceData - Commercial vehicle registration data

| LicensePlate | RegistartionId | Expired |
|--------------|----------------|---------|
| SVT 6023 | 285429838 | 1 |
| XLZ 3463 | 362715656 | 0 |
| QMZ 1273 | 876133137 | 1 |
| RIV 8632 | 992711956 | 0 |
| ... | ... | |

Events and Time

Every event that flows through the system has a timestamp

ASA supports:

- Arrival Time - Event timestamps based on arrival time (input adapter clock, e.g., Event Hubs)

- App Time - Event timestamps based on a timestamp field in the actual event tuple

User can pick up App Time from the payload

```
SELECT * FROM EntryStream TIMESTAMP BY EntryTime
```

System can assign timestamps automatically based on the event arrival time

```
SELECT * FROM EntryStream
```

Filters and Projections

From the incoming stream find only vehicles that:

- Are from either WA and CA state
- Have a weight less than 3000 lbs
- Have License plate number end in 999
- Have a make that starts with a "M"

Display:

"Passenger" if type = 1

"Commercial" if Type = 2

"Other" for all other types

Display time as 'Mins', 'Seconds', 'Milliseconds'

```
SELECT VehicleCategory =  
    Case Type  
        WHEN 1 THEN 'Passenger'  
        WHEN 2 THEN 'Commercial'  
        ELSE THEN 'Other'  
    END,  
TollId, State LicensePlate, State, Make,  
Model, Weight,  
DATEPART(mi,EntryTime) AS 'Mins',  
DATEPART(ss,EntryTime) AS 'Seconds',  
DATEPART(ms,EntryTime) AS 'Milliseconds'  
FROM EntryStream TIMESTAMP BY EntryTime  
WHERE (State = 'CA' OR State = 'WA')  
AND Weight < 3000  
AND CHARINDEX ('M', model) = 0  
AND PATINDEX('%999', LicensePlate) = 5
```

Temporal Joins

Report the time in seconds required for vehicles to pass the toll booth

```
SELECT ES.TollId, ES.EntryTime, EX.ExitTime, ES.EntryTime, ES.LicensePlate
DATEDIFF(minute, ES.EntryTime, EX.ExitTime )
FROM EntryStream ES TIMESTAMP BY EntryTime

JOIN ExitStream EX TIMESTAMP BY ExitTime

ON (EX.TollId= ES.TollId and ES.LicensePlate = EX.LicensePlate)

AND DATEDIFF(minute, ES, EX) BETWEEN 0 AND 15
```

Temporal Left Outer Join to Detect Patterns

Reports all cars that have entered the toll booth but have not exited within 5 minutes

```
SELECT ES.TollId, ES.EntryTime, ES.LicensePlate
FROM EntryStream EN TIMESTAMP BY EntryTime

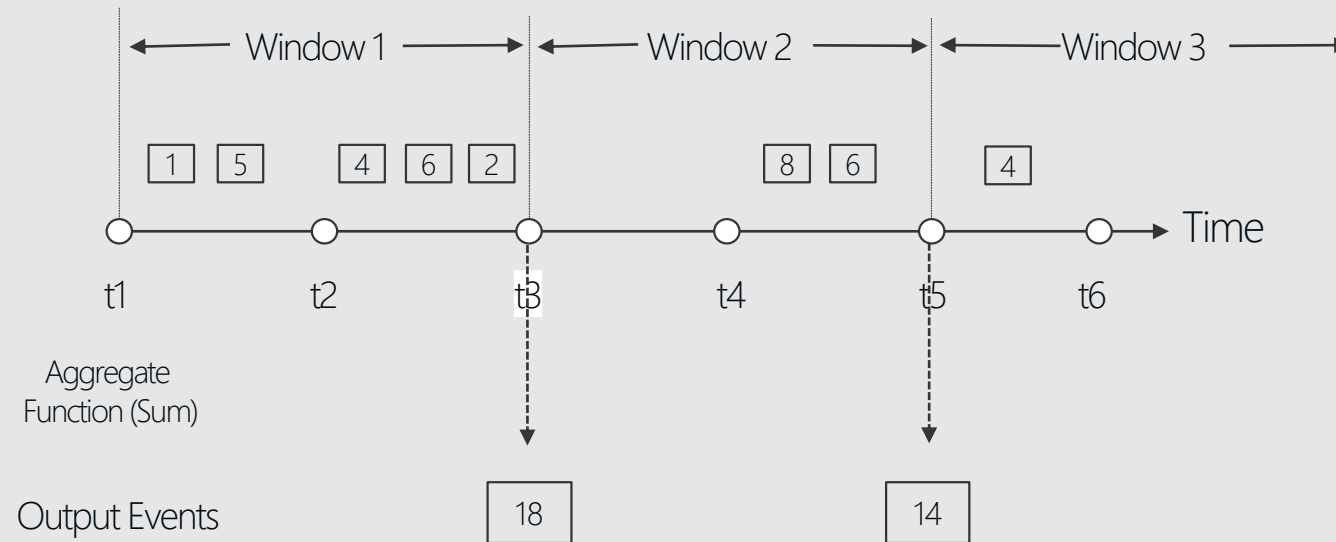
LEFT OUTER JOIN ExitStream EX TIMESTAMP BY ExitTime
ON (EN.TollId= EX.TollId AND EN.LicensePlate = EX.LicensePlate)
AND DATEDIFF(minute, EN, EX) BETWEEN 0 AND 5
WHERE EX.ExitTime IS NULL
```


Windowing Concepts

Output at the end of each window

Windows are fixed length

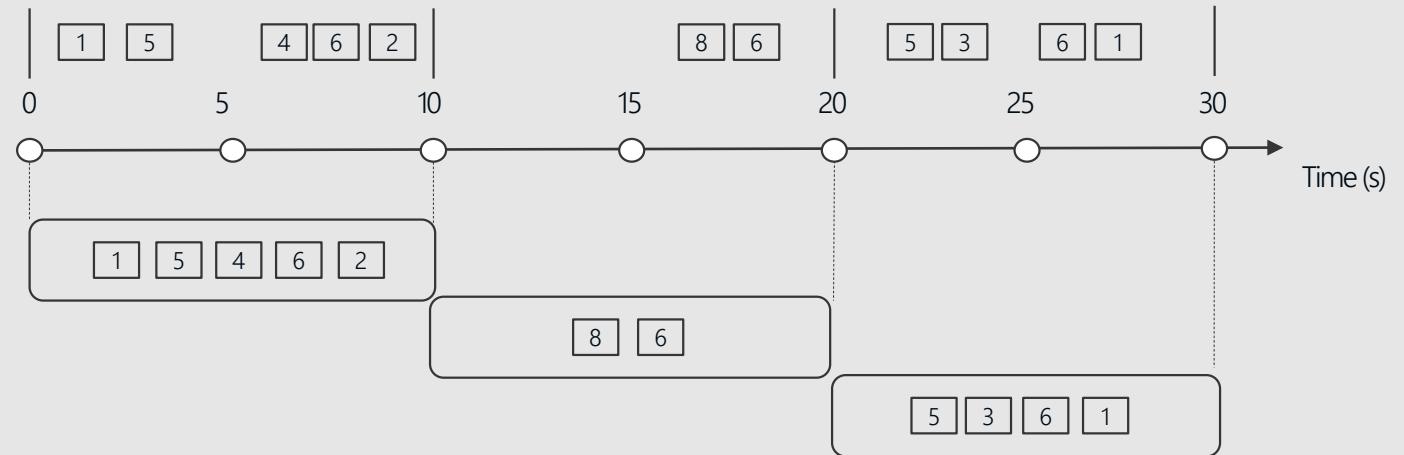
Used in a GROUP BY clause



Tumbling Windows

Every 10 seconds give me the count of vehicles entering each toll booth over the last 10 seconds

A 10-second Tumbling Window

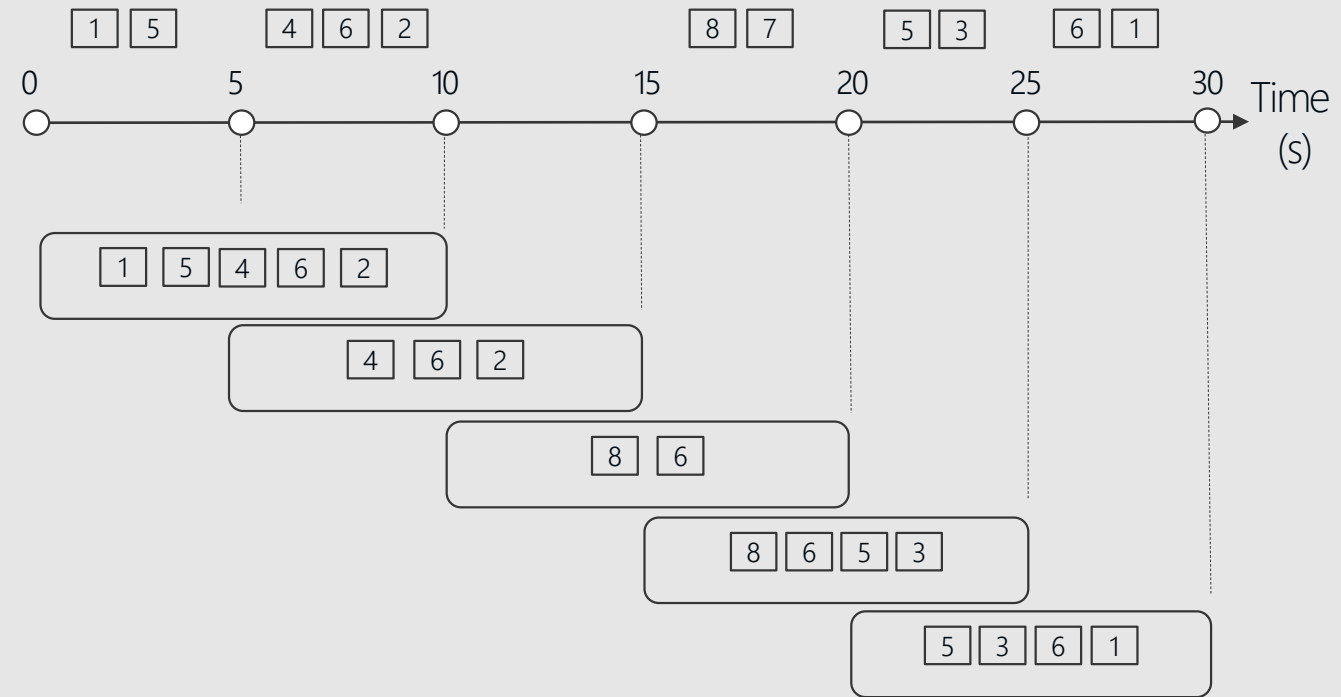


```
SELECT TollId, Count(*)  
FROM EntryStream TIMESTAMP BY EntryTime  
GROUP BY TollId, TumblingWindow(second, 10)
```

Hopping Windows

Every 5 seconds give me the count of vehicles entering each toll booth over the last 10 seconds

A 10 second Hopping Window with a 5 second hop

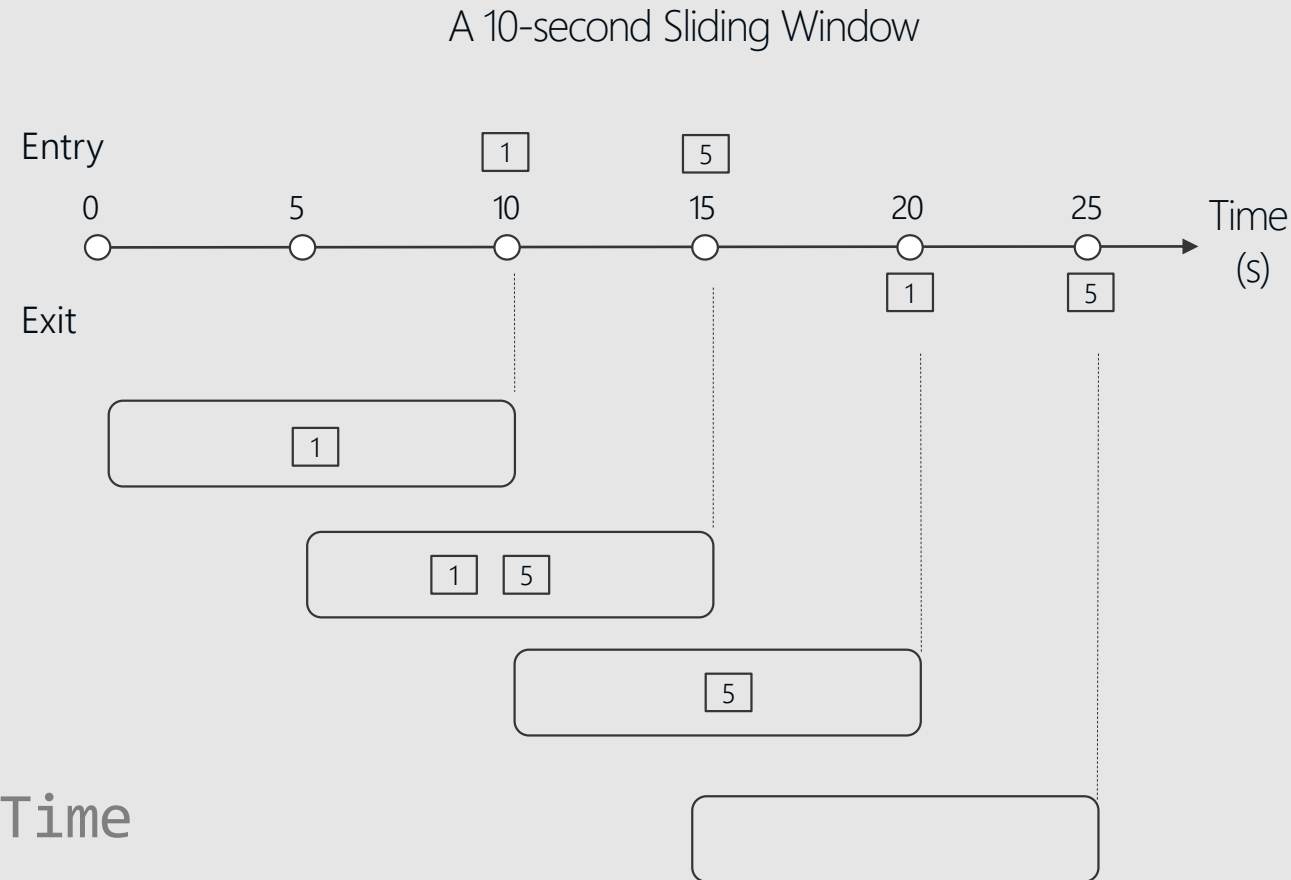


```
SELECT TollId, Count(*)  
FROM EntryStream TIMESTAMP BY EntryTime  
GROUP BY TollId, HoppingWindow(second, 10, 5)
```

Sliding Windows

Find all toll booths that have served more than 10 vehicles in the last 20 seconds

```
SELECT TollId, Count(*)  
FROM EntryStream TIMESTAMP BY EntryTime  
GROUP BY TollId, SlidingWindow(second, 20)  
HAVING Count(*) > 10
```



An output is generated whenever an event either enters/leaves the system

Advanced Query Examples



Determine if a Value has Changed

```
SELECT
    Make,
    Time
FROM
    Input TIMESTAMP BY Time
WHERE
    LAG(Make, 1) OVER (LIMIT DURATION(minute, 1)) <>
    Make
```

LAG is used to peek into the input stream one event back and get the Make value. Then compare it to the Make on the current event and output the event if they are different

Find First Event in a Window

Find first car in every 10 minute interval

```
SELECT
    LicensePlate,
    Make,
    Time
FROM
    Input TIMESTAMP BY Time
WHERE
    IsFirst(minute, 10) = 1
```


Find Last Event in a Window

Find last car in every 10 minute interval

```
WITH LastInWindow AS
(
    SELECT
        MAX(Time) AS LastEventTime
    FROM
        Input TIMESTAMP BY Time
    GROUP BY
        TumblingWindow(minute, 10)
)
SELECT
    Input.LicensePlate,
    Input.Make,
    Input.Time
FROM
    Input TIMESTAMP BY Time
    INNER JOIN LastInWindow
    ON DATEDIFF(minute, Input, LastInWindow) BETWEEN 0 AND 10
    AND Input.Time = LastInWindow.LastEventTime
```

There are two steps in the query – the first one finds latest timestamp in 10 minute windows; the second joins results of the first query with original stream to find events matching last timestamps in each window

Detect Duration of a Condition

Find out how long a condition occurred for. For example, suppose that a bug that resulted in all cars having an incorrect weight (above 20,000 pounds) – we want to compute the duration of the bug.

```
WITH SelectPreviousEvent AS
(
  SELECT
    *,
    LAG([time]) OVER (LIMIT DURATION(hour, 24)) as previousTime,
    LAG([weight]) OVER (LIMIT DURATION(hour, 24)) as previousWeight
  FROM input TIMESTAMP BY [time]
)

SELECT
  LAG(time) OVER (LIMIT DURATION(hour, 24) WHEN previousWeight < 20000
) [StartFault],
  previousTime [EndFault]
FROM SelectPreviousEvent
WHERE
  [weight] < 20000
  AND previousWeight > 20000
```

Fill Missing Values

For the stream of events that have missing values, produce a stream of events with regular intervals. For example, generate event every 5 seconds that will report the most recently seen data point.

```
SELECT System.Timestamp AS windowEnd
SELECT
  System.Timestamp AS windowEnd,
  TopOne() OVER (ORDER BY t DESC) AS lastEvent
FROM
  input TIMESTAMP BY t
GROUP BY HOPPINGWINDOW(second, 300, 5), TopOne() OVER (ORDER BY t DESC) AS
lastEvent FROM input TIMESTAMP BY t GROUP BY HOPPINGWINDOW(second, 300, 5)
```

This query will generate events every 5 second and will output the last event that was received before. Hopping Window duration determines how far back the query will look to find the latest event (300 seconds in this example).

Detect Duration Between Events

Find the duration of a given event. For example, given a web clickstream determine time spent on a feature.

```
SELECT
```

```
[user], feature, DATEDIFF(second, LAST(Time) OVER (PARTITION BY [user], feature  
LIMIT DURATION(hour, 1) WHEN Event = 'start'), Time) as duration
```

```
FROM input TIMESTAMP BY Time
```

```
WHERE
```

```
    Event = 'end'
```

Reference Data



Correlation of Event Streams with Reference Data

Static or slowly-changing data stored in blobs

Scanned for changes on a settable cadence

Joins between streams and reference data sources for correlations

Reference data appears like another input in the query

```
SELECT myRefData.Name, myStream.Value  
FROM myStream  
JOIN myRefData  
ON myStream.myKey = myRefData.myKey
```



Reference Data Example

Reports all vehicles that entered the toll with expired licenses

```
SELECT ES.EntryTime, ES.LicensePlate, ES.TollId,  
RD.RegistrationId  
FROM EntryStream ES TIMESTAMP BY EntryTime  
  
JOIN RegistrationData RD  
ON ES.LicensePlate = RD.LicensePlate  
WHERE RD.Expired = 1
```


Custom Code Support



JavaScript User Defined Functions (UDFs)

Custom Code is Supported Using JavaScript UDFs

- Stateless

- Side-effect-free

- Implementers do not need to concern themselves with sharding, resilience, or resumption

Use Cases

- ASA supports declarative representation of the logically difficult parts of the streaming computations

- Custom-code extensibility is meant for logically simple but technically difficult logic:

 - String parsing and manipulation [e.g. `Regexp_Replace()` and `Regexp_Extract()`]

 - Array operations (e.g. sorting, joining, find, fill)

 - Regular Expressions

 - Mathematics operations

 - Date operations

Restrictions

- Callouts to external REST endpoints

- Pulling reference data from an external source

- Custom event format serialization/deserialization on inputs/outputs.

- Custom aggregation functions

JavaScript UDF Example

```
SELECT
    time,
    udf.hex2Int(offset) AS IntOffset
INTO Output
FROM InputStream
```

Local definition file for the JavaScript UDF:

```
{
  "properties": {
    "type": "Scalar", //Function type. Scalar is the only supported value
    "properties": {
      "inputs": [ // Function input parameter(s).
        {
          "dataType": "any", // Input data type
        }
      ],
      "output": { // Output
        "dataType": "any" // Output data type
      },
      "binding": {
        "type": "Microsoft.StreamAnalytics/JavascriptUdf",
        "properties": { // Function definition
          "script": "function hex2Int(hexValue) {return parseInt(hexValue,
16);}",
        }
      }
    }
  }
}
```

Machine Learning Integration



Azure Machine Learning Callouts in public preview

Perform real-time scoring on streaming data


Anomaly Detection and Sentiment Analysis are common use cases

Function calls from the query

Azure ML can publish web endpoints for operationalized ML models

Azure Stream Analytics binds custom function names to such web endpoints











```
SELECT text, sentiment(text) AS score  
FROM myStream
```



Geospatial Capabilities



Real-time Geospatial Analytics Scenarios





| | | | |
|---|---|---|------------------------------------|
|  | Phone Tracking Across Cell Sites |  | Personnel Tracking & Crowd Control |
|  | Connected Car - Remote Management & Diagnostics |  | Ride Sharing |
|  | Asset Tracking |  | Geofencing |
|  | Fleet Management |  | Racecar Telemetry |
|  | Facilities Management |  | Connected Manufacturing |

and many more...

Geospatial Functions

CreatePoint
CreatePolygon
CreateLineString

ST_DISTANCE
ST_WITHIN
ST_OVERLAPS
ST_INTERSECTS



Geospatial Examples

Generate an event when gas is less than 50 km from the car

```
SELECT Cars.Location, Station.Location
FROM Cars c
JOIN Station s ON ST_DISTANCE(c.Location, s.Location) < 50 * 1000
```

Generate an event when fuel level is lower than 50%, a gas station is in promotion and course of car is pointing to gas station

```
SELECT Cars.gas, Cars.Location, Cars.Course, Station.Location,
Station.Promotion
FROM Cars c
```

```
JOIN Station s ON Cars.gas < 0.5 AND Station.Promotion AND
ST_OVERLAPS(c.Location, c.course)
```

Generate an event when a store is within a possible flooding zone

```
SELECT Store.Polygon, Flooding.Polygon
FROM Cars c
JOIN Flooding f ON ST_OVERLAPS(s.Polygon, f.Polygon)
```

Generate an event when a storm is heading mv wav

```
SELECT Cars.Location, Storm.Course
FROM Cars c
JOIN Storm s ON ST_OVERLAPS(c.Location, s.Cours
```




Combination of clustering and heat maps. Clusters are represented using color coded geometric shapes that fit together evenly



Monitoring and Troubleshooting



Job Level Monitoring



Time Range
past hour custom

Chart type
Bar Line

- ☐ Data Conversion Errors
- ☐ Failed Function Requests
- ☐ Function Events
- ☐ Function Requests
- ☐ Input Event Bytes
- ☒ Input Events
- ☒ Late Input Events
- ☒ Out of order Events
- ☒ Output Events
- ☐ Runtime Errors
- ☐ SU % Utilization

Add an alert rule

* Resource
[Dropdown]

* Name
[Text] ✓

Description
[Text] Description

* Metric
[Dropdown]

90%
85%
80%
75%
70%
65%
60%

6 PM Sep 16 6 AM 12 PM

* Condition
greater than [Dropdown]

* Threshold
90 [Text] ✓ %

* Period
Over the last 5 minutes [Dropdown]

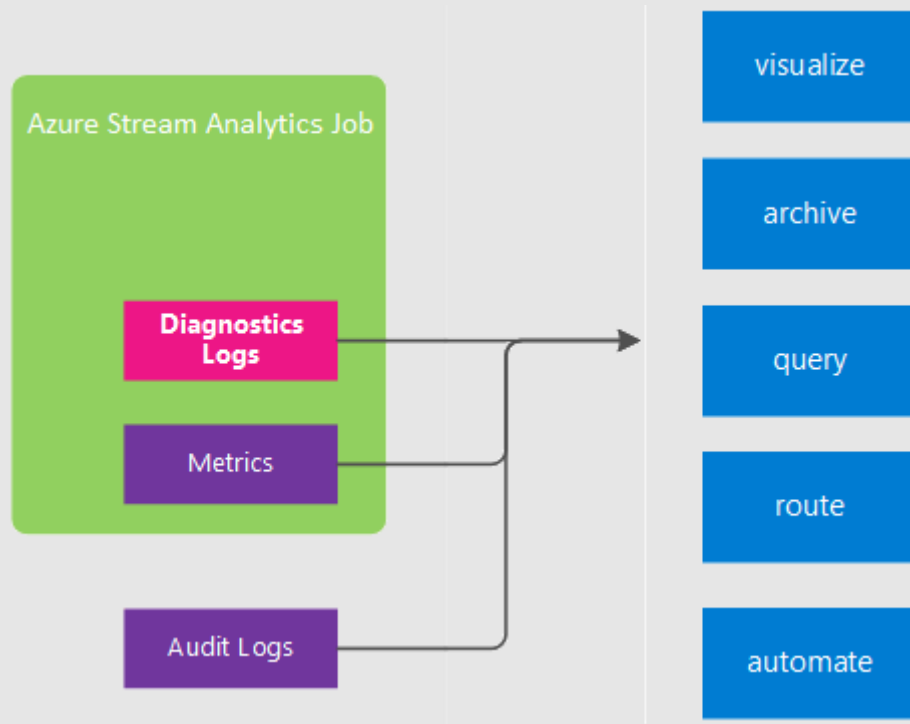
Email owners, contributors, and readers
☒

Additional administrator email(s)
admin@contoso.com [Text] ✓

Webhook
http://www.contoso.com/dowork?param [Text] ✓
[Learn more about configuring webhooks](#)

Take Action
Run a runbook from this alert >

Diagnostic Logs




“like stderr for ASA”


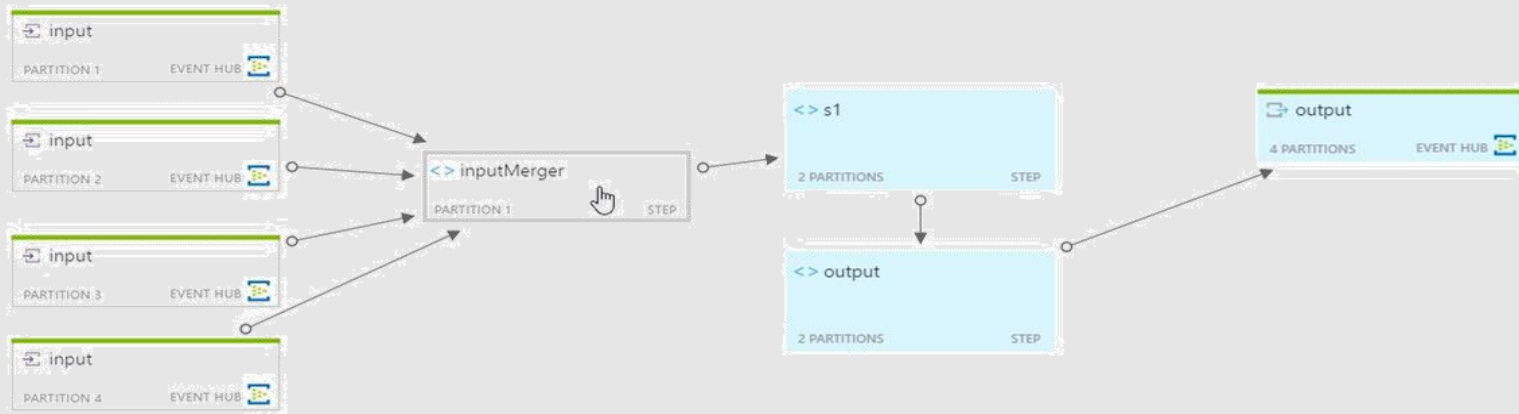
Event Hub: Monitor an ASA job with another ASA job

Storage account: Can potentially include PII data

Log Analytics

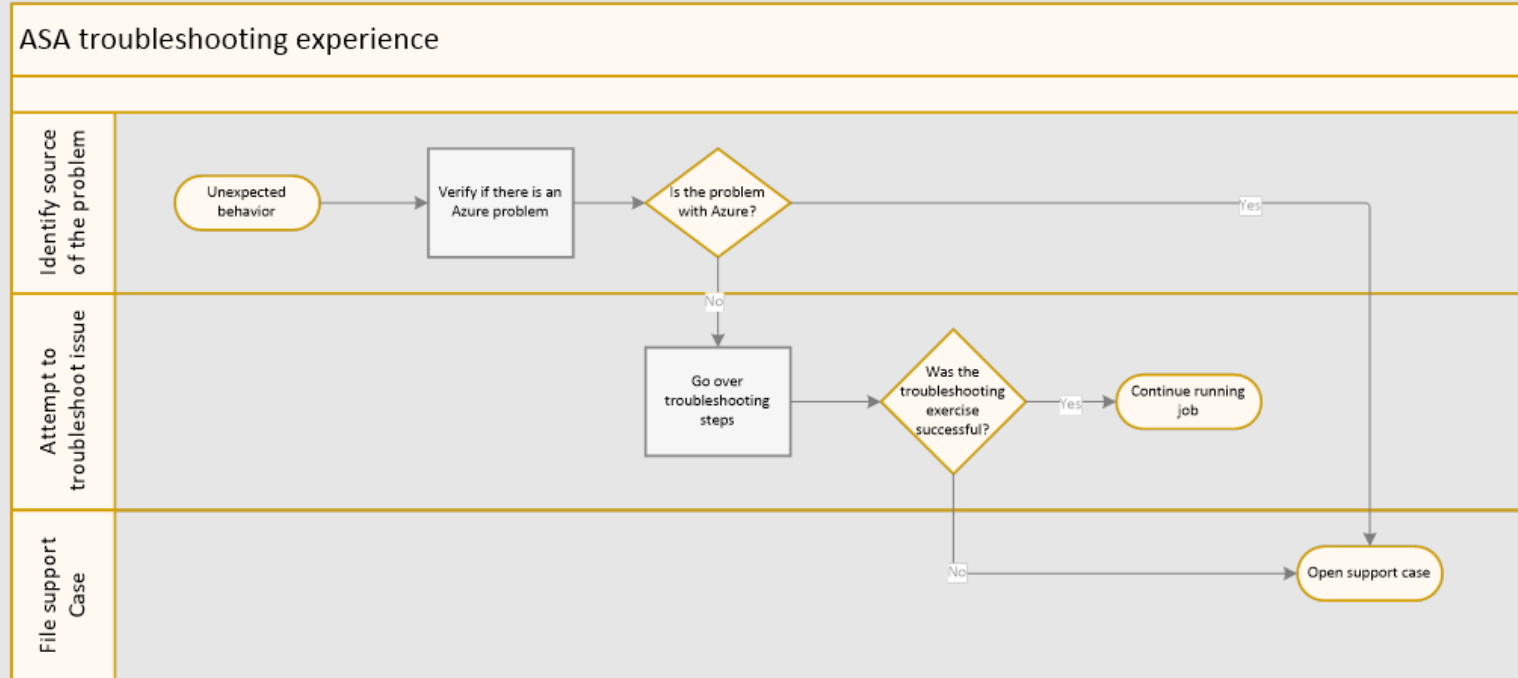


Job Diagram with Metrics



Resource Health

“is there a problem with Azure Stream Analytics that is affecting my job”



2/3, 11:18 AM

2/3, 11:41 AM

Unavailable

Your job is stopping and should stop soon

2/3, 9:19 AM

2/3, 11:18 AM

Available

There aren't any known Azure platform problems affecting this job

2/3, 9:16 AM

2/3, 9:19 AM

Unavailable

Your job may be having problems accessing its inputs or outputs


Seamless
Scalability



Streaming Units

Represents the computing resources footprint of a Azure Stream Analytics job


Seamlessly add/remove Streaming Units



Vertical Partitioning of the Query

Vertically partition into logical steps

Using a named WITH-AS in SAQL



```
WITH
StepErrors AS
(
    SELECT machine, count(*)
    FROM Errors
    GROUP BY machine
),
StepWarnings AS
(
    SELECT machine, count(*)
    FROM Warnings
    GROUP BY machine
)
SELECT machine, e.count as numErrors,
w.count as numWarnings
FROM StepErrors e JOIN StepWarnings w
ON e.machine = w.machine
WHERE e.count < w.count * 2
```

Horizontal Partitioning for Scale

Horizontal sharding into data-parallel execution over scaled-out resources

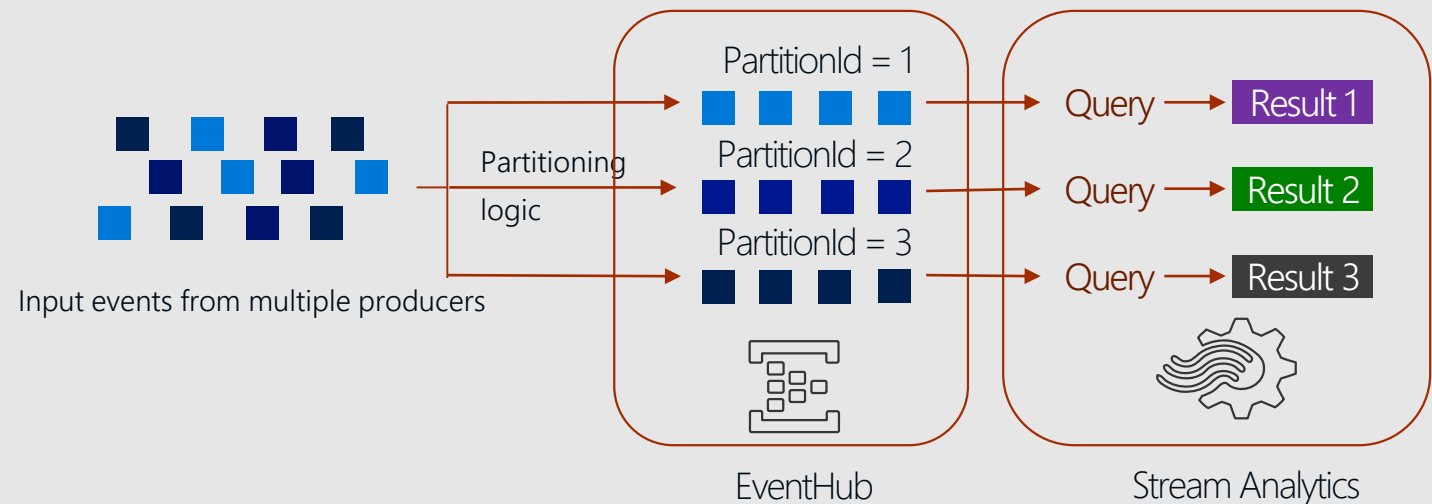
User specifies partitioning keys for each stage

Partitioning is “safe” if the partitioning key is also a grouping key
Otherwise, incomplete groupings are computed per partition

Explicit global grouping step is required to reconcile

User specifies bound on the number of SUs

Partitioning keys are assigned to shards using a built-in hash function



```
SELECT COUNT(*) AS Count, TollBoothId
FROM EntryStream PARTITION BY PartitionId
GROUP BY TumblingWindow (minute, 3), TollBoothId
```

Out-of-order and Late-arriving Events

Each individual stream is always in-order of time

Input streams that are not in-order are either:

Sorted (and therefore delayed!)

Adjusted by the system, as per a user-specified policy

Additional “punctuation” events

Advance the time in the absence of event arrivals

“Notify me when no logins occur for 3 minutes”



Tolerance for lateness

Events arrive out of order but within the tolerance: Re-ordered by timestamp

Events arrive later than tolerance: Dropped or Adjusted

Adjust - Adjusted to appear to have arrived at the latest still acceptable time

Drop - Discarded

 Save  Discard

Events that arrive late

Accept late events with a timestamp in the following range: ⓘ

Days

00 ▼

Time (HH:MM:SS)

00 ▼

00 ▼

05 ▼

Out of order events

Accept out of order events with a timestamp in the following range (MM:SS) ⓘ

00 ▼

00 ▼

Handling other events

Action ⓘ

Adjust

Drop

Interacting with Stream Analytics



Modalities of Interaction & Programmability

Azure Portal

Visual Studio

PowerShell

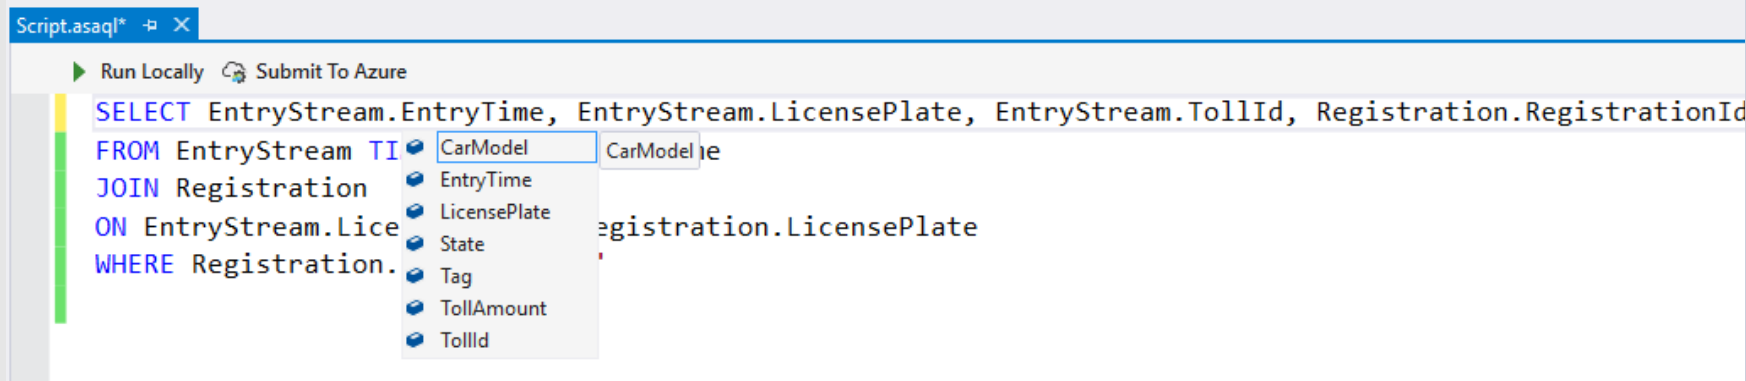
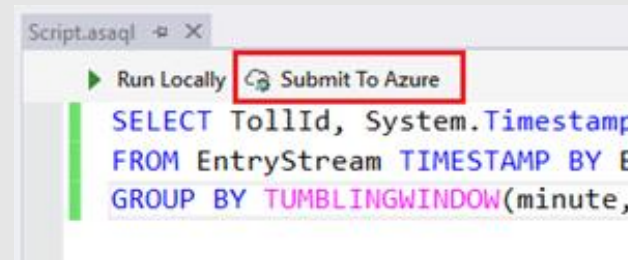
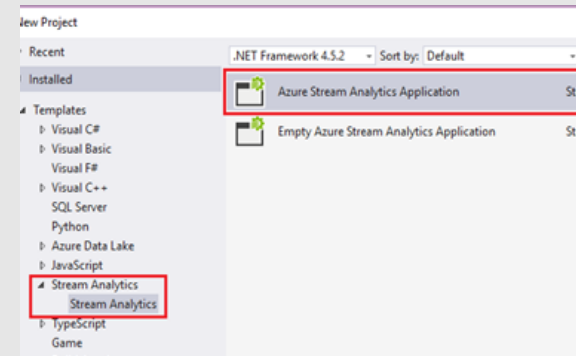
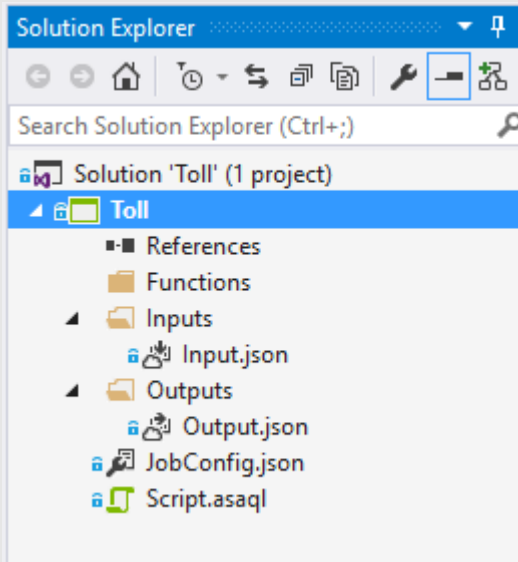
.NET SDK

REST APIs



Visual Studio Support

Version Control
Author with Intellisense
Test Locally
Deploy to Azure
Monitor



| | |
|----------------------------|------------------------|
| Job Summary | |
| Status | Running |
| Created Time | 12/12/2016 4:47:05 PM |
| Last Output Time | 12/13/2016 11:06:00 AM |
| Data Locale | en-US |
| Job Output Start Time | 12/12/2016 4:50:43 PM |
| Job Output Start Mode | CustomTime |
| Output Error Handling | Drop |
| Late Arrival Tolerance | 5 Second(s) |
| Out of Order Tolerance | 0 Second(s) |
| Out of Order Actions | Adjust |
| Stream Analytics Units | 1 |
| Job Metrics (Last 30 Mins) | |
| Total Input Events | 4,95 K |
| Input Event Bytes | 946.06 KiB |
| Total Output Events | 896 |
| Late Input Events | 11 |
| Out of Order Events | 0 |
| Function Events | 0 |
| Function Requests | 0 |
| Failed Function Requests | 0 |
| Data Conversion Errors | 0 |
| Runtime Errors | 0 |

Edge Analytics



Edge Analytics

Local Execution

Stream analytics runs on 'edge devices'

Unlock the Value of Untapped data

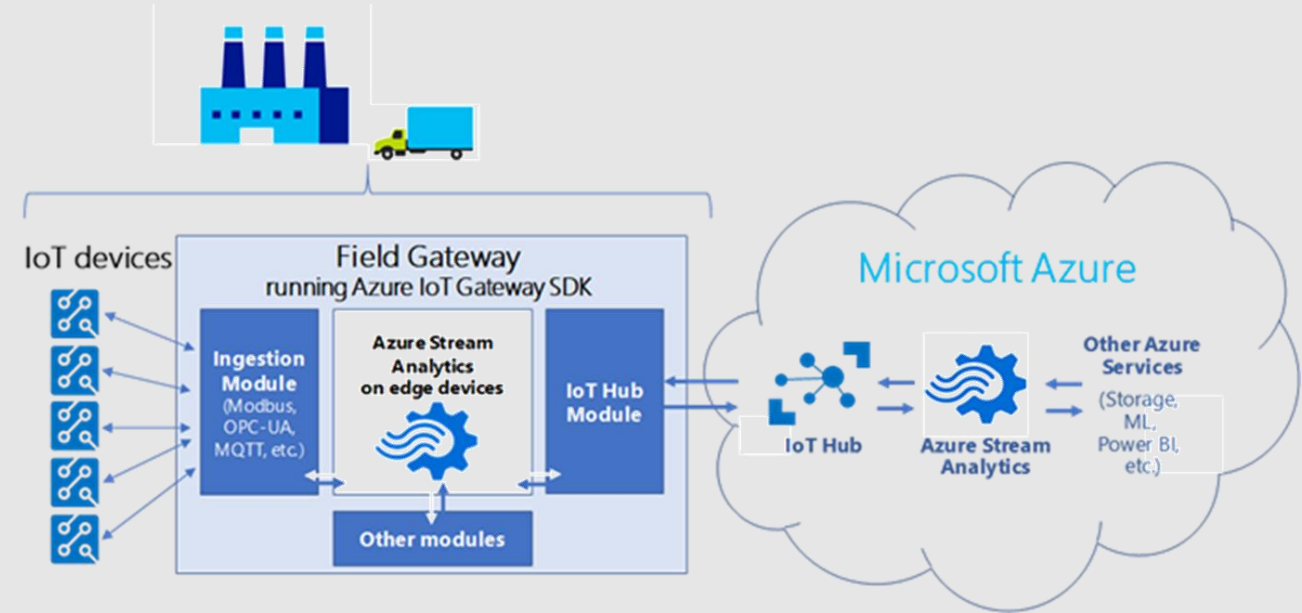
Only ~5% of data in industrial processes is sent to the cloud
Deploy intelligence near the data to unlock the full value of data

Seamless development and operations

Stream analytics jobs run in the cloud and on edge devices

Intelligent actions

Deploy situational awareness, custom code, ML models on the edge



Edge Analytics Scenarios

Low Latency

Resiliency

Efficient Use of Bandwidth

Compliance

Management at IoT-Scale

Canonical Use Cases

Reduction

When you are interested in only parts or significant changes in your operational data

Aggregation

When business operations need an aggregate view

Batching

When connectivity is intermittent and cost is high

Transformation

Converting messages from legacy industrial automation to modern applications formats

Edge Intelligence

Machine Learning models on the edge

Reference data

Edge Analytics Platform Requirements

Windows or Linux

Azure IoT Gateway

Memory

- Minified stream analytics engine needs ~2MB of main memory

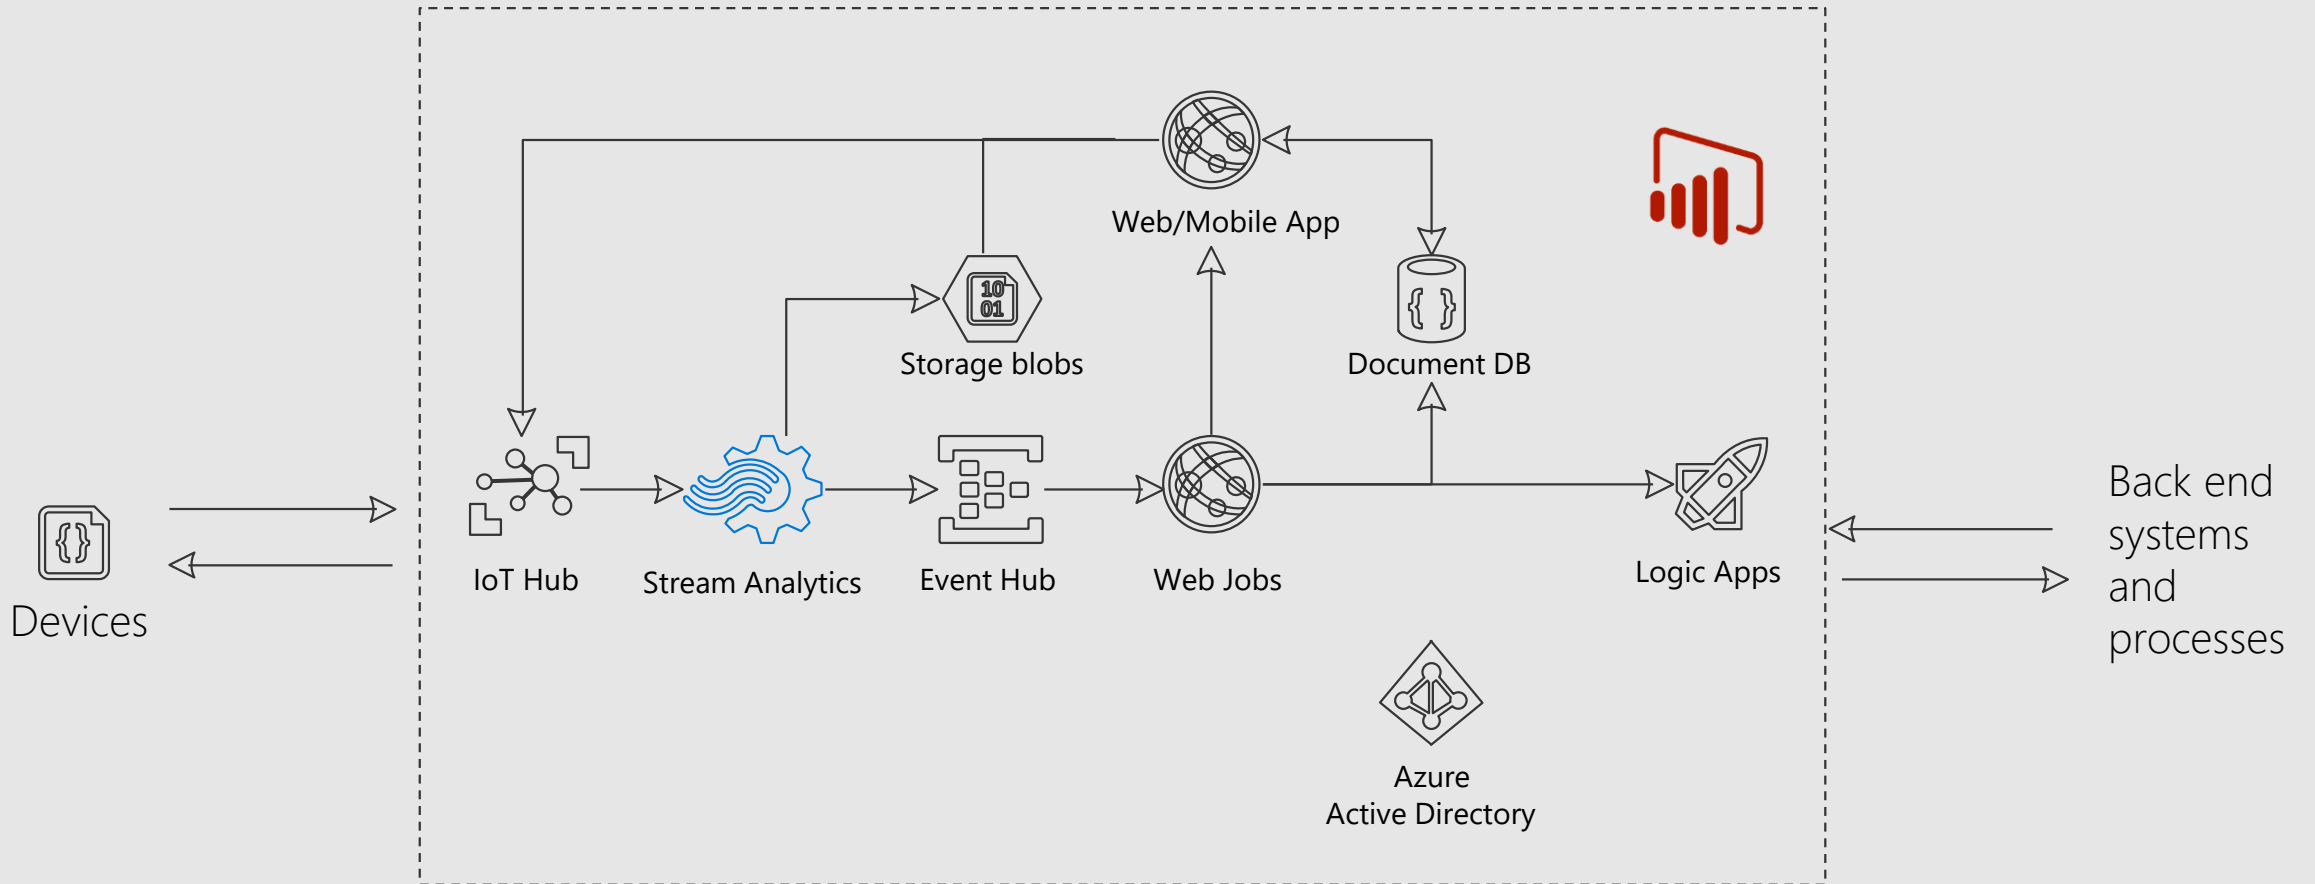
- Storage and additional memory based on the amount of data

Node.js

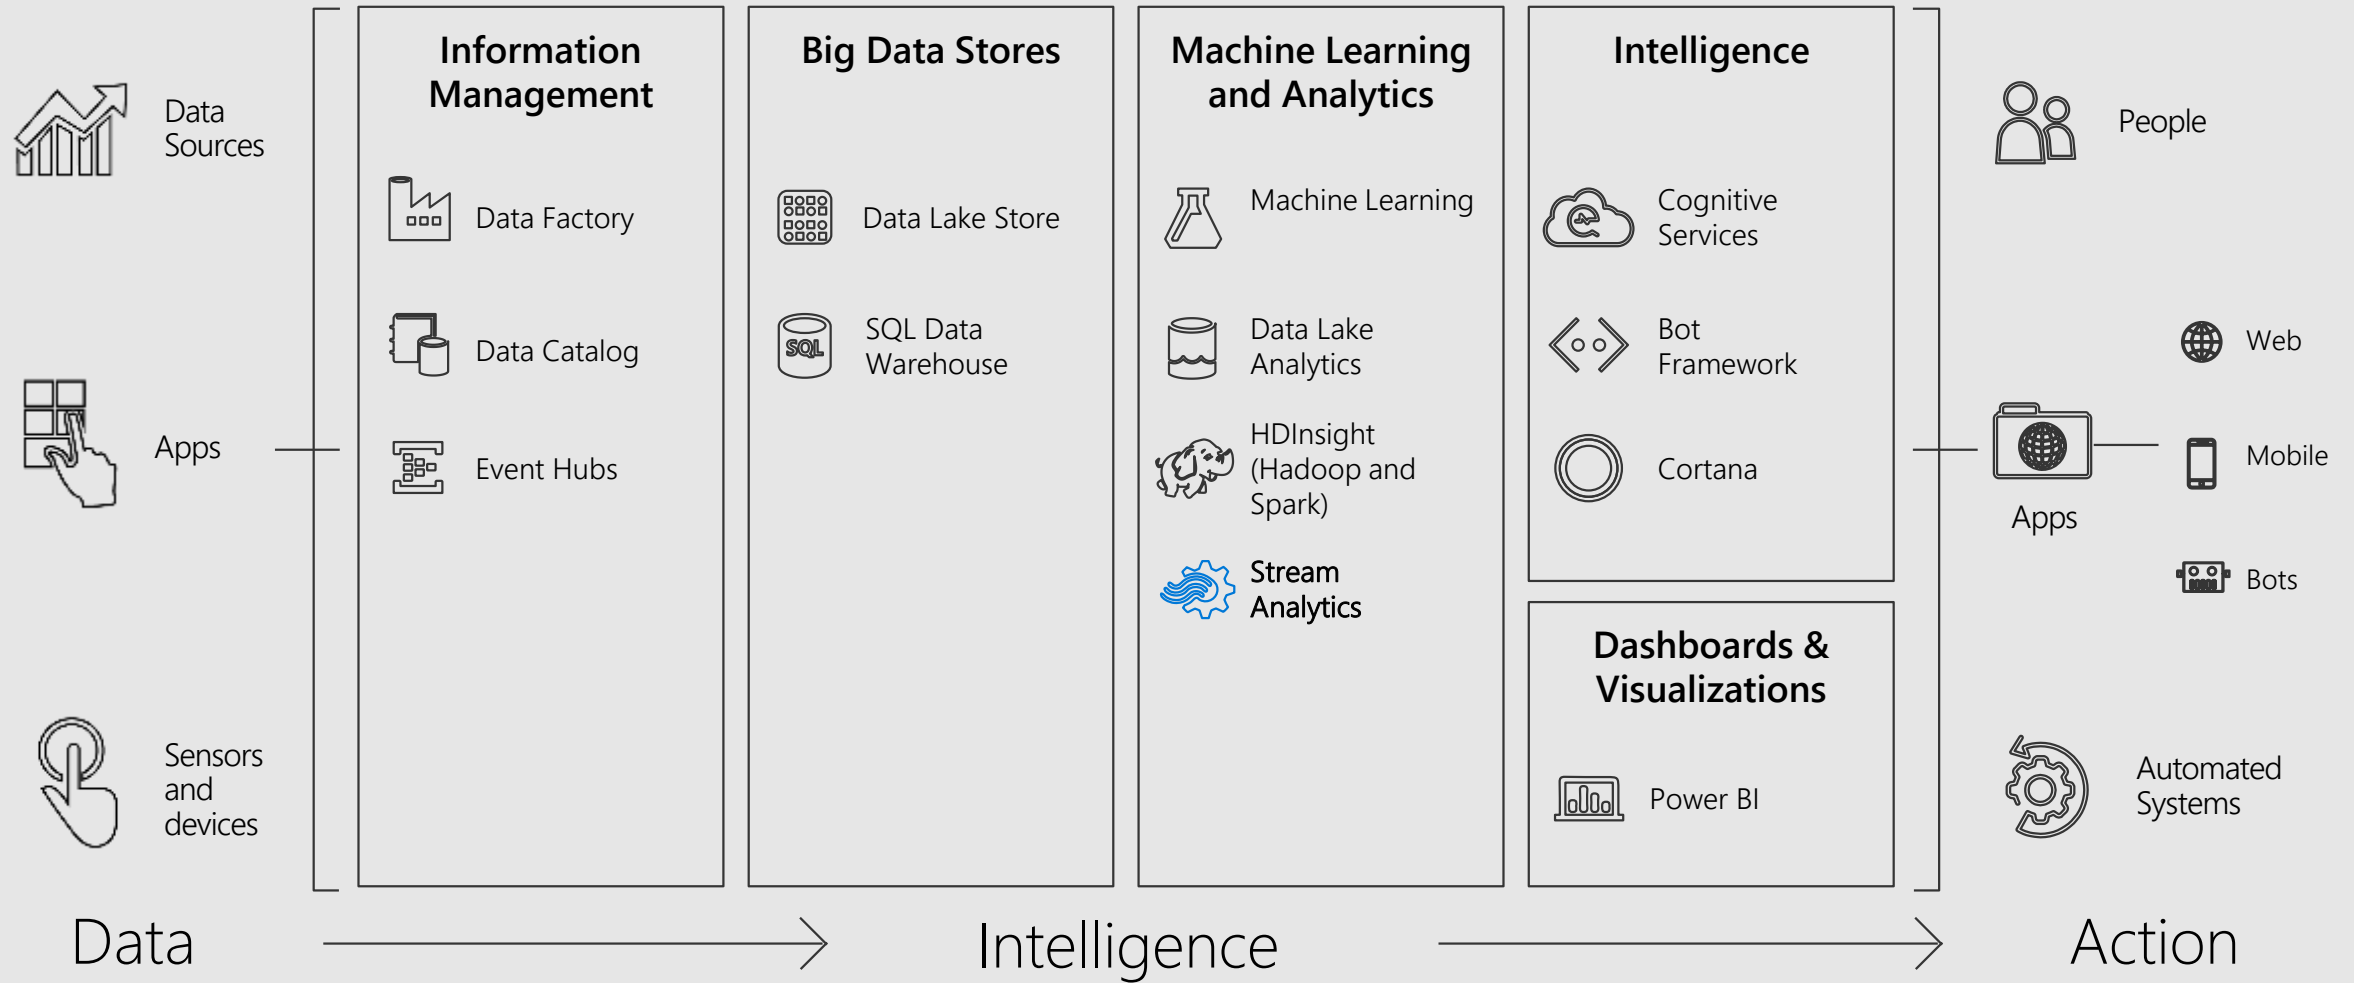
Stream Analytics is a Key Part of Multiple Strategic Bets



Microsoft IoT Suite



Microsoft Cortana Intelligence Suite



Streaming Options on Azure



Streaming Options with Microsoft

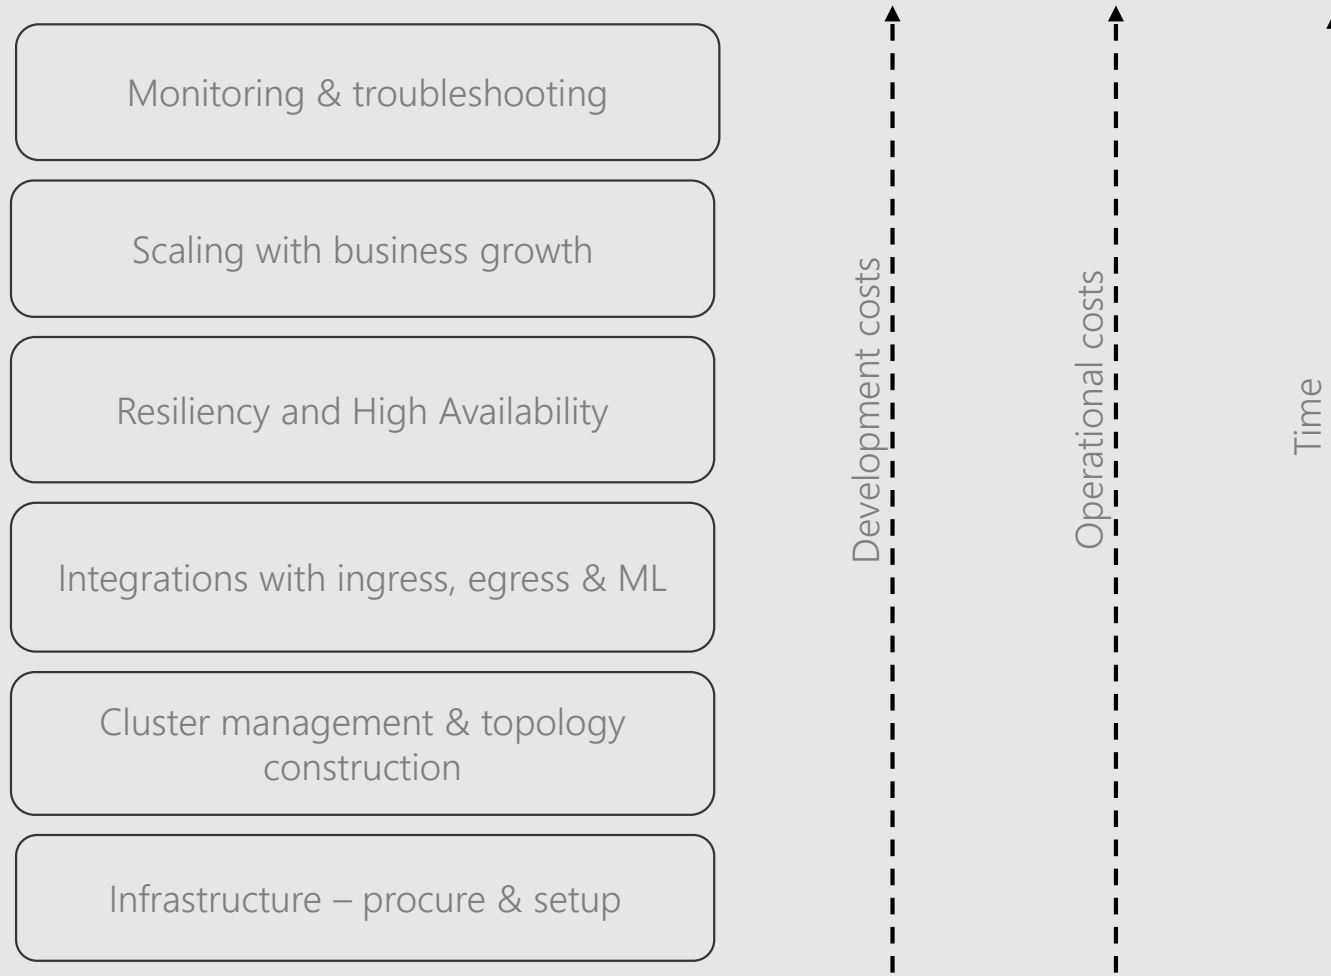
Azure Stream Analytics

Storm on Azure HDInsight

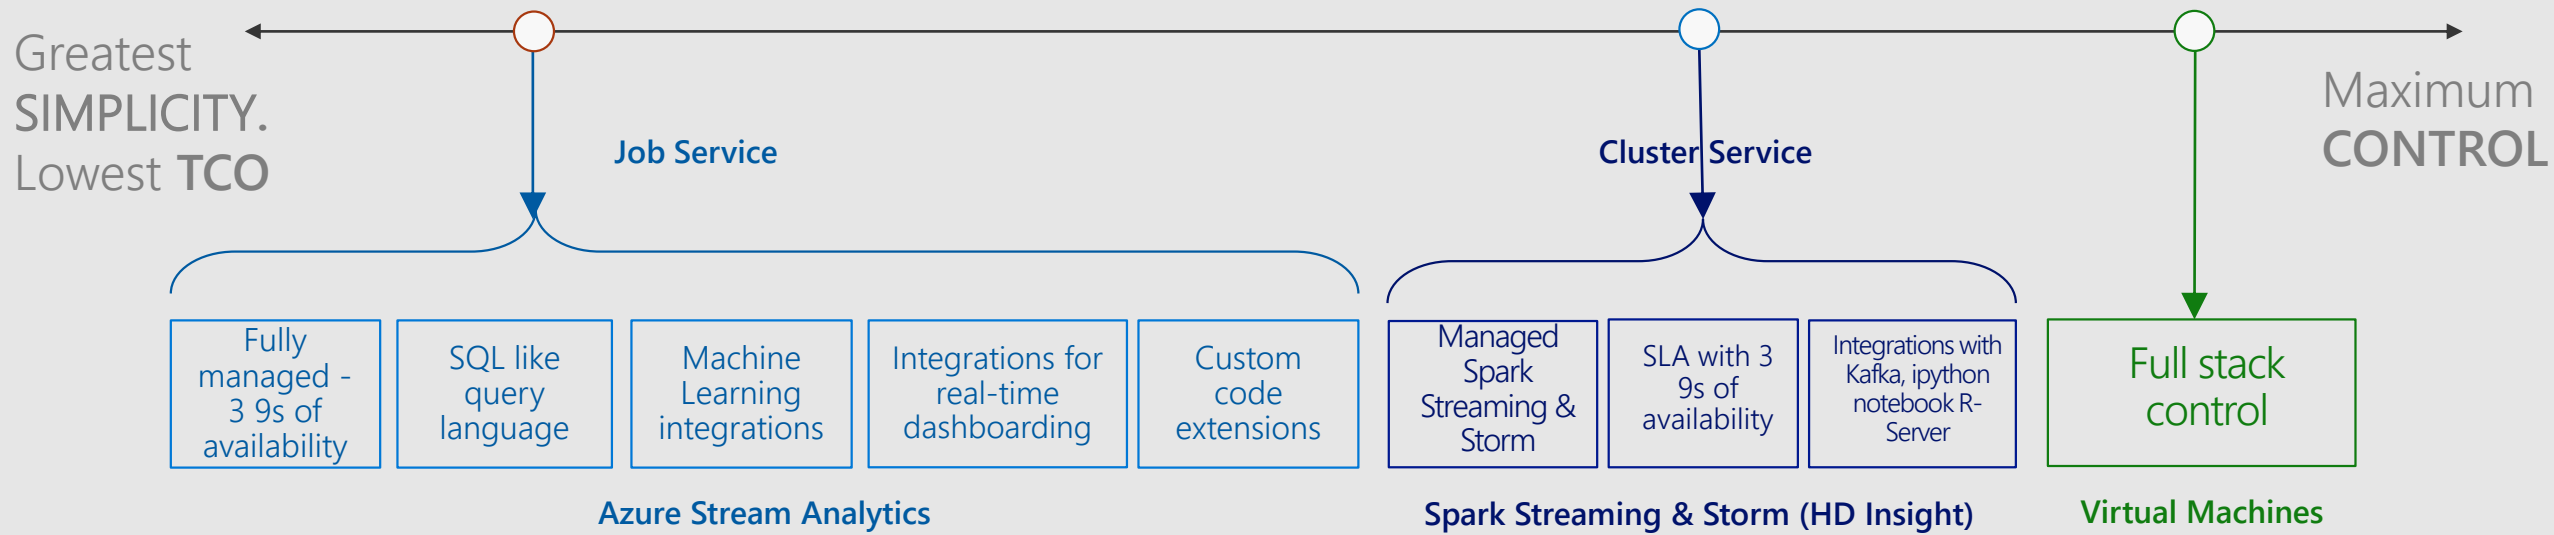
Spark Streaming on Azure HDInsight

StreamInsight in SQL Server

Streaming Stack



Tradeoff Spectrum



Apache Storm on Azure HDInsight

Open source Apache Storm as a managed service

SLA of 99.9% up time

Built in Scale-up & Scale-down capability

Customers can scale up and scale down a running cluster; with no impact to a running topology

Deep integration with Event hub

Program your Storm Topology in Java or C#

Rich Visual Studio experience

Debug, monitor, and troubleshoot

Manage and deploy topologies

Fully integrated Azure portal experience

Write your data into SQL Azure, DocumentDB, PowerBI, ...

Spark on Azure HDInsight

Managed Service

Fully supported by Microsoft and Hortonworks

Latest open source Spark 2.0 with 100+ stability fixes (available later this week on 9/30)

Enterprise Readiness

99.9% uptime SLA

Compliance: PCI, ISO 27018, SOC, HIPAA, EU-MC

Azure Integrations

Jupyter Notebooks (Scala, Python, Automatic data visualizations)

IntelliJ and Eclipse plugins (job submission, remote debugging)

ODBC connectors: Power BI, Tableau, Qlik, SAP, Excel, etc.

Technical Comparison

| | Storm on HDInsight | Spark Streaming on HD Insight | Azure Stream Analytics |
|-----------------------------|---|-------------------------------|--|
| Strictest Guarantee | At-least-once (exactly once with Trident) | Exactly once | At-least-once |
| Processing Model | Event-at-a-time (micro batching with Trident) | Micro batching | Adaptive batching (event-at-a-time) |
| Programing Language Support | Java, C# | Scala, Python, Java | SAQL: SQL like query language, JavaScript UDFs |
| Open Source | Yes | Yes | No |